

Betting and Belief: Prediction Markets and Attribution of Climate Change

John J. Nay

School of Engineering
Vanderbilt University
PMB 351826
2301 Vanderbilt Place
Nashville, TN 37235-1826, USA

Martin Van der Linden

Department of Economics
Vanderbilt University
PMB 351819
2301 Vanderbilt PL
Nashville, TN 37235-1819, USA

Jonathan M. Gilligan

Department of Earth & Environmental Sciences
Vanderbilt University
PMB 351805
2301 Vanderbilt PL
Nashville, TN 37235-1805, USA

ABSTRACT

Despite overwhelming scientific evidence, a large fraction of the American public doubts that greenhouse gases are causing global warming. We present a simulation model as a computational test-bed for climate prediction markets. Traders adapt their beliefs about future temperatures based on the profits of other traders in their social network. We simulate two alternative climate futures, in which global temperatures are primarily driven either by carbon dioxide or by solar irradiance. These represent, respectively, the scientific consensus and a hypothesis advanced by prominent skeptics. We conduct sensitivity analyses to determine how a variety of factors describing both the market and the physical climate may affect traders' beliefs about the cause of global climate change. Market participation causes most traders to converge quickly toward believing the "true" climate model, suggesting that a climate market could be useful at building public consensus.

1 INTRODUCTION

The climate change debate has become strongly polarized over the past two decades. Although the scientific consensus on the anthropogenic nature of climate change strongly increased, beliefs about climate change did not evolve much within the public (Vandenbergh et al. 2014). In addition, the divide on anthropogenic climate change between liberals and conservatives has grown steadily as the question is becoming increasingly politicized and potentially disconnected from scientific evidence. The costs of misinformed climate policies are high. If climate change is not human-induced but is believed to be so, public resources will be spent on unnecessary efforts. On the other hand, if climate change is human-induced but not recognized as so, the costs of inaction could be devastating. Given that effective climate policies require long-lasting measures to be taken, the importance of an accurate consensus on the issue is manifest.

Attempts to foster such consensus face many social and psychological challenges, some of which could be addressed by creating climate prediction markets where participants can "put their money where their mouths are" (Hsu 2011, Vandenbergh et al. 2014). The idea of using prediction markets to efficiently

aggregate information about uncertain event outcomes has been widely discussed (Horn et al. 2014). Prediction markets have interesting theoretical properties (Set and Selten 1998, Hanson 2012), and perform well in terms of prediction accuracy and information aggregation in experiments (Hanson et al. 2006, Healy et al. 2010) simulation models (Klingert and Meyer 2012, Jumadinova and Dasgupta 2011) and the real-world (Wolfers and Zitzewitz 2006, Pathak et al. , Dreber et al. 2015).

Bloch, Annan and Bowles Bloch et al. (2010) have proposed using derivatives markets to reduce the scientific uncertainty in estimating the impact of future climate change. Existing prediction markets (e.g. hypermind.com, betfair.com, and predictit.org) focus on near term events such as elections months away, so it is difficult to extrapolate empirical findings to the climate case. Furthermore, we are interested in investigating the unobservable beliefs of traders. Therefore, we turn to simulation modeling informed by climate and economic theory. We simulate a prediction market where traders exchange securities related to climate outcomes to explore whether, and under what social and climate conditions, prediction markets may be useful for increasing convergence of climate beliefs. Our work can be extended as part of a computational design process for effective climate prediction markets.

From a public policy perspective, changing the explanatory models of market participants is one of the most important roles that prediction markets might play. This is perhaps the most important social benefit of prediction markets given the comparable predictive power of statistical supervised learning models, which can provide information at a much lower cost than creating and maintaining a market (Goel et al. 2010). Effective climate policy not only requires an accurate consensus on future climate *outcomes*. It also requires an accurate consensus on the causal *mechanisms* influencing such outcomes. If people agree that temperature will rise, but some believe it will be due to greenhouse gases, while others believe that it will be caused by increased solar activity, inconsistent and ineffective policies may be implemented.

2 RELATED WORK

To the best of our knowledge, the idea that prediction markets can generate consensus *on the factors* affecting uncertain events has never been explored quantitatively. Agent-based simulations prediction markets have been studied (Klingert and Meyer 2012, Tseng et al. 2010, Jumadinova and Dasgupta 2011), including some that feature communication between agents. In these models, however, beliefs about the uncertain outcomes are constructed in rather abstract ways. In particular, beliefs are not based on structural models from which agents could derive causal implications, so these models are not suitable for investigating the convergence of the underlying explanatory models agents employ for prediction.

Tseng et al. (2010) created an agent-based model (ABM) of a continuous double auction market with multiple market strategies, including two variants of a zero-intelligence agent. They compared the behavior of their simulation with data from a prediction market for the outcome of political elections. They found that despite their simplicity, zero-intelligence agents capture some salient features of real market data. Klingert and Meyer (2012) compared the predictive accuracy different kinds of simulated markets (continuous double auctions and logarithmic market scoring rules) and reached similar conclusions to experimental work by Hanson et al. (2006). None of these models featured agents learning and updating their beliefs.

Jumadinova and Dasgupta (2011) created a continuous double auction model in which agents update their beliefs about uncertain events based on newly-acquired information. The better the information set, the more likely the agent was to put higher weight on last period's prices when revising their beliefs. However, little structure was imposed on the information set and the way it was used to generate the weights for updating beliefs, so that model does not permit studying the convergence of trader beliefs about predictive models. Ontañón and Plaza (2009) studied the effect of deliberation on a simulated prediction market, in which agents used case-based reasoning (Aamodt and Plaza 1994) to debate uncertain outcomes with their neighbors in a social network. This model could in principle be used to assess convergence of predictive models, but this was not done and the process of forming and revising beliefs through argumentation is not well suited to studying stochastic time series, such as those relevant to climate prediction markets.

3 MODEL DESIGN

Traders bet on global temperature anomalies six years in the future. During the six-year period, traders buy and sell futures. Every year, traders update their models and forecast of future temperatures based on newly available data. At the end of each six-year period, winners collect gains and traders revise beliefs about climate models, based on their ideology and the beliefs of top earners in their social network. In this section, we describe the models used to generate future temperature data, agents beliefs about those models, the market procedures, and the social network connecting agents. We conclude with details on model dynamics and an overview of the model parameters we experimentally vary.

3.1 Temperature Models

For climate time-series, we use the annual anomaly of global mean temperature. For years from 1880–2014 we use the GISTEMP global mean land-sea annual temperature anomalies (GISTEMP Team 2016, Hansen et al. 2010) and for years from 2015 onward, we project future climates under two alternative theories: In both theories, changes in global temperature are proportional to changes in a deterministic forcing plus a stochastic noise term. For simplicity, we choose two alternative expressions for the deterministic climate forcing: one, which corresponds to conventional climate science, takes the natural logarithm of the atmospheric carbon dioxide concentration in parts per million (Archer 2012)[p. 37], and the other, which corresponds to an alternative theory advocated by many who doubt or reject conventional climate science, takes the total solar irradiance (the brightness of sunlight, in Watts per square meter, at the top of the atmosphere) averaged over the 11-year sunspot cycle (Soon 2005). Most scientific models of the earth’s climate include many forcings, including carbon dioxide, other greenhouse gases, aerosols, total solar irradiance, and more. In these models, the changing CO₂ concentration is, by a large margin, the strongest forcing (IPCC 2013)[p. 14], so for simplicity we take this as the model. Choosing only one forcing term for each of the competing models simplifies comparison because each has the same number of adjustable parameters.

For CO₂ we used historical emissions through 2005, harmonized with RCP 8.5 representative concentration pathway from 2005 onward (Kolp and Riahi 2009, Riahi et al. 2011) and for TSI we used the harmonized historical values, with a projection through 2100 from Velasco Herrera et al. (2015), which seems to be the best TSI model.

Warming coefficients for each model were determined by linear regression of historical temperatures from 1880–2014 against the historical values for each model’s forcing term. The noise model was determined by fitting an ARMA(p, q) model to the residuals from the regression, using the [Stan](#) probabilistic modeling language and the [rstan](#) package (Carpenter et al. 2016, Stan Development Team 2016). We identified the optimal model for the autocorrelated noise term by performing the regression analysis for all combinations of $p, q \in [0, 1, 2]$ and using the Watanabe-Akaike Information Criterion to select the optimal noise model (Vehtari et al. 2015). In both cases (TSI and logCO₂), the optimal noise model was AR(1).

Future climates were generated by applying the warming coefficients to Eq. 1:

$$T_{\text{model}}(t) = \beta_{\text{model}} F_{\text{model}}(t) + \varepsilon, \quad (1)$$

where $T_{\text{model}}(t)$ is the temperature at time t , under a given model of what causes warming, $F_{\text{model}}(t)$ is the forcing (TSI or logCO₂) at time t , β_{model} is the regression coefficient, and ε is an ARMA(p, q) noise term. The coefficient β and the parameters of the ARMA noise model were fit to the historical data (1880–2014), as described above. An example realization of future climates for both the ln(CO₂) and TSI models is shown in Fig. 1.

3.2 Climate Beliefs

Traders use one of the two models (temperature depends on CO₂ or TSI) to forecast future temperature. These models are interpreted as the trader’s beliefs about the true climate process, the driving factor of




Figure 1: Historical measurements of temperature and a realization of possible future temperatures under two alternate models of climate physics.

long-term global temperatures. They represent pervasive positions on climate change in the public debate. In order to approximatively match the current configuration of beliefs in climate change in the United States, the CO_2 model is randomly assigned to half of the traders, while the TSI model is assigned to the other half of the traders.

Thus, both when the true data generating process is CO_2 and TSI, approximately half of the traders use the true data-generating model to make predictions. Traders using the true model do not necessarily make perfectly accurate predictions however. Although these traders believe in the correct *functional form* of the model, they still need to calibrate their model based on limited noisy data. Therefore, the values these traders assign to the parameters of the model will typically be different from the exact parameters in the data-generating process.

3.3 Traders and Markets

Traders are initially endowed with a single experimental currency unit (ECU). Traders use their model to forecast the distribution of future temperature and determine their reservation price for different securities. Each security pays 1 ECU at the end of the trading sequence if the temperature at the end of the sequence falls into a certain range.

Traders are risk-neutral expected utility maximizers. Therefore, their reservation price for a security is simply their assessment of the probability that the temperature will fall in the range covered by the security at the end of the sequence. At each time-step (year), the agents use the new year's temperature data to re-estimate the coefficients for the model they believe explains climate change, using Bayesian linear regression with an AR(1) noise model, and use the joint posterior probability distribution of regression and noise coefficients to estimate the probability distribution for the temperature at the end of the current trading sequence. Stan was sufficiently fast that we could perform a full Bayesian analysis at each time step. Traders then use this posterior probability distribution to assign reservation prices for buying and selling securities.

Based on their reservation price, agents behave as “zero-intelligence” traders (Gode and Sunder 1993). They attempt to sell securities at a random price above their reservation, and to buy securities at a random




Figure 2: An example of the social network among traders. Yellow and blue circles represent traders who believe in the two different models of climate change. Lines indicate social network connections. In a highly segmented network (seg close to 1), most links connect like-minded traders and few connect traders with opposing beliefs.

price below their reservation. These trading strategies are simple but provide accurate approximations of behavior in prediction markets (Klingert and Meyer 2012), and in financial markets more broadly.

Based on traders' sell and buy orders, traders exchange securities following a continuous-double auctions (CDA) procedure (see Model Dynamics below for more details). CDAs or some close variants are common procedures to match buy and sell orders. CDA are notably used on large stock markets (Tseng et al. 2010).

3.4 Social Network

Traders are part of a social network where each agent forms two links at random, and then forms links randomly, ensuring that each agent is connected to at least two other agents. Every time securities are realized, each trader looks at the performance of her richest neighbor in the network. Traders start with the same initial amount of ECU and differences in ECU can only come from market interactions. Therefore, if some trader is poorer than her richest neighbor, the trader interprets it as a signal that her richest neighbor may have a better model of the climate. Then, the trader considers adopting the model of her richest neighbor. For each trader, the willingness to revise her belief is determined by how ideologically loaded her belief is, which is a parameter we vary in our experiments.

An example of a snapshot of a social network is depicted in Fig. 2. We assume that the initial network is segregated: traders who initially believe in CO_2 forcing (resp. TSI) are more likely to be linked with other traders who also believe in CO_2 (resp. TSI). This reflects reality where cultural and social groups are highly correlated with climate change beliefs. Although traders can change their approximate model over time, the connections between traders do not change as the market unfolds, i.e the edges are fixed.

3.5 Model Dynamics

The time periods t are grouped into trading *sequences*. In a given sequence, the potential payments associated with traded securities are all based on the temperature at the end of the sequence. For instance,

the third trading sequence might start in period $t = 1964$ and end in period $t = 1970$. In this case, a security traded in the third sequence pays 1 ECU if the temperature at $t = 1970$ falls into the range of temperatures covered by the security.

At each time t , traders are assumed to know the past value of the temperature $T_{0:t}$, carbon dioxide $\text{CO2}_{0:t}$, and total solar irradiance $\text{TSI}_{0:t}$. In a sequence finishing at time t^* , traders also have common knowledge of $\text{CO2}_{t:t^*}$ and $\text{TSI}_{t:t^*}$, the future values of carbon dioxide and total solar irradiance up to t^* . However, at any t , traders do not know the value of any future temperatures. Thus, the traders know the forcing terms for the rest of the trading sequence, but do not know the value of T_{t^*} . Traders can only predict T_{t^*} using their approximate model and their knowledge of $T_{0:t}$, $\text{CO2}_{0:t^*}$, and $\text{TSI}_{0:t^*}$. Notice that because $T_{0:t}$, $\text{CO2}_{0:t^*}$, and $\text{TSI}_{0:t^*}$ are common knowledge, in each period t , any two trader with the same approximate model m forms the same stochastic beliefs about future temperatures $P_{t,m}(T_{t^*} | T_{0:t}, \text{CO2}_{0:t^*}, \text{TSI}_{0:t^*})$. The probability distribution $P_{t,m}$ incorporates both epistemic uncertainty (the trader does not know the true values of the coefficients of the climate model) and aleatory uncertainty (in addition to deterministic warming, the temperature exhibits stochastic noise).

At each time t , traders:

1. recalibrate their approximate model based on the new set of temperature data available at t ,
2. use the posterior probability distribution from (1) to assign beliefs about the probability distribution of future temperatures at time t^* : $P_{t,m}(T_{t^*} | T_{0:t}, \text{CO2}_{0:t^*}, \text{TSI}_{0:t^*})$ and use it to determine the expected value they attach to each security, and
3. trade on the CDA market as follows:
 - Every trader i chooses at random a security s_i^B she will try to buy.
 - Every trader i also chooses at random a security s_i^S she will try to sell among the securities she owns a positive amount of (if any).
 - Traders then decide of their selling price p_i^S and buying price p_i^B . To do so, traders first compute their expected values $E(s_i^B)$ and $E(s_i^S)$ for securities s_i^B and s_i^S . Then traders set p_i^S at random above $E(s_i^S)$ and p_i^B at random below $E(s_i^B)$ (see Model Parameters below for more details).
 - Traders go to the market one at the time, in an order drawn randomly for each t .
 - When trader i comes to the market, she places limit orders in the order book. These orders specify that i is willing to buy s_i^B at any price below p_i^B , and to sell s_i^S at any price above p_i^S .
 - The market maker then tries to match i 's orders with some order which was put in the book before i came to the market.
 - If there are outstanding sell offers for s_i^B at a price below p_i^B , a trade is concluded. Trader i buys one unit of s_i^B from the sellers who sells at the lowest price below p_i^B , and the sell and buy offers are removed from the order book.
 - If there are outstanding buy offers for s_i^S at price above p_i^S , a trade is concluded. Trader i sells one unit to the buyer who buys at the highest price above p_i^S , and the sell and buy offers are removed from the order book.
 - Whenever all traders have come to the market, any remaining outstanding offer is removed from the order book, and the trading period is concluded.

At t^* , when the sequence ends, there is only one security s^* associated with a range of temperatures including the actual temperature T_{t^*} . At t^* , traders:

1. receive 1 ECU per unit of s^* they own, and
2. consider adapting their neighbors' approximate model as described in Social Network.

3.6 Model Parameters

The model depends on the following parameters, which we vary in simulation experiments to determine their effects on the convergence of beliefs. We group the parameters into climate, network and individual behavioral factors.

- Climate parameters:
 - `true.model`: temperature data-generating process (CO₂ or TSI).
- Network parameters:
 - `n.traders`: the number of traders.
 - `n.edg`: the number of edges in the social network.
 - `seg`: the initial degree of homophily in the network. The higher `seg`, the higher the initial homophily.
- Behavioral parameters:
 - `risk.tak`: determines the distribution of risk tolerance with respect to successfully buying or selling securities. Higher risk tolerance corresponds to demanding more aggressive prices for buying and selling, and hence, a higher risk of not consummating a trade that would be mutually advantageous to the buyer and seller.
 For each trader i , the level of `risk.taki` is drawn uniformly at random from $[0, \text{risk.tak}]$. The higher `risk.taki`, the higher the price i will demand for selling securities, and the lower the price i will offer for buying.
 Formally, let `reservi,t` be the reservation price of trader i at time t for some security s . Then trader i picks her buying or selling price for s uniformly at random in the interval $[(1 - \text{risk.tak}_i)\text{reserv}_{i,t}, \text{reserv}_{i,t}]$ for buying and $[\text{reserv}_{i,t}, (1 + \text{risk.tak}_i)\text{reserv}_{i,t}]$ for selling.
 - `ideo`: determines the degree of “ideology” of traders. For each trader i , the level of `ideoi` is drawn uniformly at random from $[0, \text{ideo}]$. If `ideo` is high, traders will not revise their approximate models easily, even when faced with evidence that their richest neighbor is doing better than them. Formally, for each trader i and each sequence, a parameter d_i is drawn from $[0, \text{ideo}_i]$. The value of d_i is the probability that i adopts the approximate model of her richest neighbor if that neighbor is doing better than i at the end of the sequence (in monetary terms).

4 RESULTS

4.1 Historical Climate

As a validation that the market is operating correctly, we run the market with actual temperatures 1880 to 2015 and market betting from 1931 to 2014. Before 1980 there is no statistically significant warming trend so we would not expect the market to change beliefs before then. As a global warming signal begins to show up in 1980s, traders begin to converge toward the $\log(\text{CO}_2)$ model (Fig. 4).

4.2 Future Scenarios

Our primary focus is on simulating the market in future scenarios. This is more relevant for policy design and more interesting theoretically due to the increasing divergence of global temperature values under the two future scenarios. The sensitivity analysis is based on a Latin hypercube sampling of parameter sets from the following distributions (Beachkofski and Grandhi 2002, Carnell 2012).

- `ideo` \sim Uniform(0,1)
- `market.complet` \sim Uniform(0,1000) (mapped into integer)
- `n.edge` \sim Uniform(100,200) (mapped into integer)
- `n.traders` \sim Uniform(50,250) (mapped into integer)




Figure 3: Estimated effects of model parameters on convergence of beliefs in future scenario. Positive numbers accelerate convergence and negative numbers retard it. True model refers to the true climate model being CO₂, so traders converge to the true model more quickly when carbon dioxide drives climate change than when total solar irradiance does. **TODO: John: explain the error bars: one sigma, two sigma, 95% CI, ...**

- `risk.tak` \sim Uniform(0,1)
- `seg` \sim Uniform(0,1)
- `true.model` \sim Bernoulli(0.5)

The future CO₂ forcing is taken from the RCP 8.5 scenario (Riahi et al. 2011) and the future TSI forcing is taken from a prediction of 21st Total Solar Irradiance by Velasco Herrera et al. (2015). Because of the stochastic noise term, the temperature time series were different in each simulation.

We use the model to perform 10,000 simulations for each input paramter set, and then conducted a partial rank correlation coefficient analysis on the relationship between the input matrix, X , and the resulting simulated outcome vector of belief convergence scores, y (Marino et al. 2008, Pujol et al. 2014, Saltelli et al. 2009). Partial correlation computes the linear relationship between the part of the variation of X_i and y that are linearly independent of other X_j ($j \neq i$). The only difference between the partial correlation and the partial rank correlation that we use here is that the outcome variable y is first ranked-transformed in order to allow us capture potentially non-linear relationships: every value, y_s , in the data set is replaced by a number $f(y_s) \equiv |\{y_r \mid y_s > y_r\}|$, where y_r are the observed values of y in the data set.

The sensitivity analysis averages over time and thus masks time trends. We randomly drew from the above distributions for ideology, risk tolerance, and the number of traders, and conducted an experiment crossing 95th and 5th percentile values for the number of edges per trader and the segmentation of the social network and both values for the true model. We then collected the time series of convergence across these eight designs for visualization of the distributions of convergence over time (Fig. 4). Under most parameterizations, within 15–20 years, the median fraction of traders believing in the true model is over 75%.

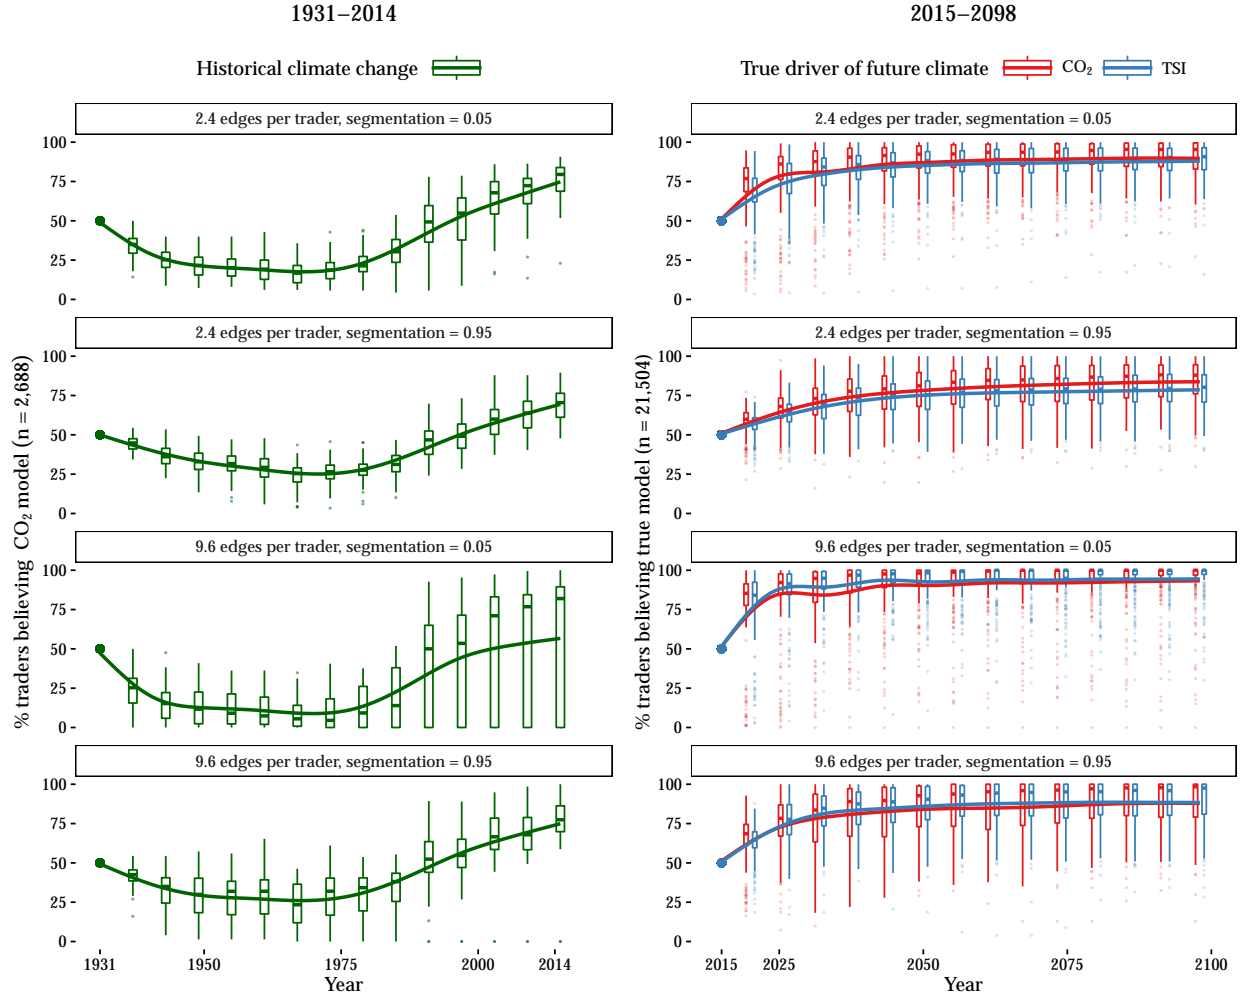


Figure 4: Convergence over trading sequences for different degrees of social-network connection and segmentation.

5 DISCUSSION

Our sensitivity analysis, Fig. 3, shows that ideology, risk tolerance, the number of traders, the number of edges per trader, the segmentation of the social network, and the true model all statistically significantly affect convergence. Increasing the number of edges per trader increases the flow of information through the market, which causes traders to converge toward believing the “true” climate model.

Segmented social networks reduce convergence by creating an “echo-chamber” effect, in which lack of interaction between traders with different views reduces traders’ access to information that could persuade them to change their beliefs. This is apparent in Fig. 4, where the rate of convergence in highly segmented markets ($seg = 0.95$) is considerably slower, especially in the first three trading sequences (18 years), than for markets with low segmentation ($seg = 0.05$).

When the true model is CO_2 -driven there is more convergence toward the true model. We believe this is because when models are fit to the historical data, the residuals from the CO_2 -driven model are considerably smaller than those from the TSI-driven model. Thus, projections of future climate change using the TSI model will exhibit significantly larger stochastic noise than the CO_2 models, and this noise makes it more difficult for traders to identify the true model when the true model is TSI.

The higher the risk tolerance, the greater the range of prices at which a trader will buy and sell securities. This creates greater opportunities for traders to profit if their model is correct and to lose money if it is incorrect. This increased volatility puts more information in play, and we believe this is why greater risk taking enables traders to identify the correct model more quickly. We have observed a similar phenomenon in a very different context, where adding stochastic noise to player decisions in iterated games improves the accuracy with which machine-learning algorithms can identify the players' strategies (Nay and Gilligan 2015).

What is not clear to us is why increasing the number of traders slows convergence, even if the number of edges per trader remains fixed. This will be a topic for future research.

The historical sequence in Fig. 4 shows convergence to the TSI model until the early 1970s, after which traders converge toward the CO₂ model. This is consistent with the fact that greenhouse gas concentrations did not become high enough to begin dominating natural climate variations until the 1970s or so, and it took until the 1990s for the temperature record to show clear signs of anthropogenic interference (IPCC 2013).

6 CONCLUSIONS

We simulate two alternative climate futures: one where the natural log of CO₂ is the primary driver of global temperature and one where sun spots are the primary driver. These represent the two most plausible competing views in the public discourse and our analysis is agnostic about which is "true".

Market participation causes traders to converge toward believing the "true" climate model under a variety of model parameterizations in a relatively short time: In markets with low segmentation, the number of traders believing the true model rises from 50% to 75% in roughly 12 years, and even in highly segmented markets, 75% convergence is achieved in 18–24 years. This suggests that a climate prediction market could be useful for producing broad agreement about the causes of climate change, and could have persuasive power for people who are not persuaded by the overwhelming consensus among climate researchers that greenhouse gases (and especially CO₂) are responsible for the majority of observed climate change (IPCC 2013). However, market segmentation has a large effect on the speed of convergence, so transparency and effective communication about the performance of traders with different beliefs will be important.

The fact that rapid convergence to the true model occurs regardless of which model is actually "true" may persuade those who doubt the scientific consensus that the market is ideologically neutral, and that the deck is not stacked to produce a pre-determined result.

All code for the model is available at johnjnay.com/predictMarket. This project is a computational test-bed for public policy design: our code can be extended to test the effects that trading strategies, cognitive models, future climate scenarios, and market designs have on the evolution of trader beliefs.

ACKNOWLEDGEMENTS

We thank Yevgeniy Vorobeychik for feedback on this project. We thank Manuel Naumovich Velasco for sharing the data from his Total Solar Irradiance prediction. We also thank the authors of [R](#) (R Core Team 2014) and [Stan](#) (Carpenter et al. 2016). This manuscript was prepared using knitr (Xie 2014).

REFERENCES

- Aamodt, A., and E. Plaza. 1994. "Case-Based Reasoning: Foundational Issues, Methodological Variations, and System Approaches". *AI Communications* 7:39–59.
- Archer, D. 2012. *Global Warming: Understanding the Forecast*. 2nd ed. Wiley.
- Beachkofski, B., and R. Grandhi. 2002. "Improved Distributed Hypercube Sampling". In *43rd AIAA/ASME/ASCE/AHS/ASC Structures, Structural Dynamics, and Materials Conference*. American Institute of Aeronautics and Astronautics.

- Bloch, D., J. Annan, and J. Bowles. 2010. "Cracking the Climate Change Conundrum with Derivatives". *Wilmott Journal* 2:271–287.
- Carnell, R. 2012. "lhs Latin Hypercube Samples". R package, Comprehensive R Archive Network.
- Carpenter, B., A. Gelman, M. Hoffman, D. Lee, B. Goodrich, M. Betancourt, M. A. Brubaker, J. Guo, P. Li, and A. Riddell. 2016. "Stan: A probabilistic programming language". *Journal of Statistical Software*. (in press).
- Dreber, A., T. Pfeiffer, J. Almenberg, S. Isaksson, B. Wilson, Y. Chen, B. A. Nosek, and M. Johannesson. 2015. "Using prediction markets to estimate the reproducibility of scientific research". *Proceedings of the National Academy of Sciences* 112:15343–15347.
- GISTEMP Team 2016. "GISS Surface Temperature Analysis (GISTEMP)". Technical report, NASA Goddard Institute for Space Studies.
- Gode, D. K., and S. Sunder. 1993. "Allocative Efficiency of Markets with Zero-Intelligence Traders: Market as a Partial Substitute for Individual Rationality". *Journal of Political Economy* 101:119–137.
- Goel, S., D. M. Reeves, D. J. Watts, and D. M. Pennock. 2010. "Prediction Without Markets". In *Proceedings of the 11th ACM Conference on Electronic Commerce, EC '10*, 357–366. New York, NY, USA: ACM.
- Hansen, J., R. Ruedy, M. Sato, and K. Lo. 2010. "Global surface temperature change". *Reviews of Geophysics* 48:RG4004.
- Hanson, R. 2012. "Logarithmic Market Scoring Rules for Modular Combinatorial Information Aggregation". *The Journal of Prediction Markets*:3–15.
- Hanson, R., R. Oprea, and D. Porter. 2006. "Information aggregation and manipulation in an experimental market". *Journal of Economic Behavior & Organization* 60:449–459.
- Healy, P. J., S. Linardi, J. R. Lowery, and J. O. Ledyard. 2010. "Prediction Markets: Alternative Mechanisms for Complex Environments with Few Traders". *Management Science* 56:1977–1996.
- Horn, C. F., B. S. Ivens, M. Ohneberg, and A. Brem. 2014. "Prediction markets—A literature review 2014". *The Journal of Prediction Markets* 8:89–126.
- Hsu, S.-L. 2011. "A Prediction Market for Climate Outcomes". *University of Colorado Law Review* 83:179–256.
- IPCC 2013. *Climate Change 2013: The Physical Science Basis*. Cambridge University Press.
- Jumadinova, J., and P. Dasgupta. 2011. "A multi-agent system for analyzing the effect of information on prediction markets". *International Journal of Intelligent Systems* 26:383–409.
- Klingert, F. M., and M. Meyer. 2012. "Comparing Prediction Market Mechanisms Using An Experiment-Based Multi-Agent Simulation." In *Proceedings 26th European Conference on Modelling and Simulation*, edited by K. G. Troitzsch, Möhring, and U. Lotzmann, 654–661.
- Kolp, P., and K. Riahi. 2009. "RCP Database". Technical report.
- Marino, S., I. B. Hogue, C. J. Ray, and D. E. Kirschner. 2008. "A methodology for performing global uncertainty and sensitivity analysis in systems biology". *Journal of Theoretical Biology* 254:178–196.
- Nay, J. J., and J. M. Gilligan. 2015. "Data-Driven Dynamic Decision Models". In *Proceedings of the 2015 Winter Simulation Conference*, edited by L. Yilmaz, W. Chan, I. Moon, T. Roeder, C. Macal, and M. Rosetti, 2752–2763: IEEE Press.
- Ontañón, S., and E. Plaza. 2009. "Argumentation-Based Information Exchange in Prediction Markets". In *Argumentation in Multi-Agent Systems*, edited by I. Rahwan and P. Moraitis, Number 5384 in Lecture Notes in Computer Science, 181–196. Springer Berlin Heidelberg.
- Pathak, D., D. Rothschild, and M. Dudik. "A comparison of forecasting methods: Fundamentals, polling, prediction markets, and experts". *Journal of Prediction Markets* 9:1–31.
- Pujol, G., B. Iooss, and A. J. Lemaitre. 2014. "sensitivity: Sensitivity Analysis". R package, Comprehensive R Archive Network.
- R Core Team 2014. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing.

- Riahi, K., S. Rao, V. Krey, C. Cho, V. Chirkov, G. Fischer, G. Kindermann, N. Nakicenovic, and P. Rafaj. 2011. "RCP 8.5A scenario of comparatively high greenhouse gas emissions". *Climatic Change* 109:3357.
- Saltelli, A., K. Chan, and E. M. Scott. 2009. *Sensitivity Analysis*. Chichester: Wiley.
- Set, R., and R. Selten. 1998. "Axiomatic characterization of the quadratic scoring rule". *Experimental Economics* 1:43–62.
- Soon, W. W.-H. 2005. "Variable solar irradiance as a plausible agent for multidecadal variations in the Arctic-wide surface air temperature record of the past 130 years". *Geophysical Research Letters* 32:L16712.
- Stan Development Team 2016. "RStan: The R interface to Stan". R package, Comprehensive R Archive Network.
- Tseng, J.-J., C.-H. Lin, C.-T. Lin, S.-C. Wang, and S.-P. Li. 2010. "Statistical properties of agent-based models in markets with continuous double auction mechanism". *Physica A: Statistical Mechanics and its Applications* 389:1699–1707.
- Vandenbergh, M. P., K. T. Raimi, and J. M. Gilligan. 2014. "Energy and Climate Change: A Climate Prediction Market". *UCLA Law Review* 61:1962–2017.
- Vehtari, A., A. Gelman, and J. Gabry. 2015. "Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC". arXiv:1507.04544.
- Velasco Herrera, V. M., B. Mendoza, and G. Velasco Herrera. 2015. "Reconstruction and prediction of the total solar irradiance: From the Medieval Warm Period to the 21st century". *New Astronomy* 34:221233.
- Wolfers, J., and E. Zitzewitz. 2006. "Prediction markets in theory and practice". Working Paper 12083, NBER.
- Xie, Y. 2014. *Dynamic Documents with R and knitr*. Boca Raton: Chapman & Hall/CRC.

AUTHOR BIOGRAPHIES

JOHN J. NAY is a Ph.D. Candidate at Vanderbilt University. He develops and applies machine learning and simulation methods to better predict behavior in complex systems. His research involves collaboration with social scientists, environmental scientists, computer scientists and engineers. His email address is john.j.nay@vanderbilt.edu, and website: <http://johnjnay.com/>.

MARTIN VAN DER LINDEN is a Ph.D. Candidate at Vanderbilt University.

TODO: Martin, please fill this in.

His email address is martin.van.der.linden@vanderbilt.edu, and website: <https://martinvdlinden.wordpress.com/>.

JONATHAN M. GILLIGAN is Associate Professor of Earth and Environmental Sciences at Vanderbilt University and Associate Director for Research at the Vanderbilt Climate Change Research Network. His research and teaching interests center on integrating natural science, social science, behavioral science, and engineering to study coupled human-natural systems. He holds a Ph.D. in Physics from Yale University. His email address is jonathan.gilligan@vanderbilt.edu, and website: <https://my.vanderbilt.edu/jonathangilligan/>.